


Fig 2. Sparse MLP / expert dataflow (one graph-capturable launch chain)



MoE: routed rows sorted by expert; each column-block carries its expert id, indexing stacked 2:4-sparse weights (routed-row scheduling, not per-expert launches)