

Fig 6. Distributed sparse-FP4 MoE (DeepSeek-V4-Flash, expert-parallel, no NVLink)

Compute vs communication

Expert-parallel scaling

